

# Performance Analysis and Optimization Report

This section shows the main bottlenecks that we had encountered during the development of our StudyBuddy database and website in phase 3.

## Jacob Craig

### 1. Data Loading

- In phase II we switched the clean data used for course\_resources which allowed for more visually appealing data in the database for the resource table
- This new data allowed for the second version of CourseResourceCleaner
  - Normalizes text via functions strip\_weird and normalize snippet
  - Writes a cleaned CSV

### 2. Query Optimization

- SearchCourseSmart: Originally was doing full table scans on courses
  - Added a CREATE INDEX idx\_courses\_code\_name
  - MySQL range scans on course\_code and course\_name rather than the entire table
  - Complexity improved from O(N) to O(Log N + k) where k is the the number of matching rows
- Group Membership & Join Requests: Membership had no constraints regarding making owners, ensuring that users could not join a group if the max capacity was reached, dealing with multiple users trying to join at the same time, etc. Duplicates on StartDirectConversation started occurring a lot in tests, there needed to be a way to stop duplicate direct messages.
  - Added a CREATE INDEX idx\_gm\_group\_user
  - Added a CREATE INDEX idx\_jr\_group\_user\_status
  - Procedures such as JoinGroupWithLock, KickGroupMember, RequestJoinPublicGroup, and ApproveJoinRequest now use index lookups instead of scanning by group
  - Membership tests ( WHERE group\_id = ? AND user\_id = ? ) are effectively O(log N) scans instead of O(N) scans as the number of members/requests grow
- Direct Messaging + Inbox: loading direct message conversations and messages required sorting direct\_message tables by time. If the direct\_messages table became large, this would be slow.
  - Added a CREATE INDEX idx\_dm\_convo\_time

- Added a CREATE INDEX idx\_dc\_users
- GetDirectMessages and GetDmInboxForUser tables now do range scans on conversation\_id, sent\_at instead of sorting in mem
- StartDirectConversation table now can check if a conversation already exists using idx\_dc\_users, stopping duplicates from happening with an easy index lookup

### 3. Matching Algorithm

- The first version of the StudyBuddy Match tried to get compatibility with study groups. Since the StudyBuddy schema cannot tell exactly what course a user is in, the original used study groups to see a users course. This made it impossible for a user to be matched with another user without joining a study group.
  - Made it such that students can select up to five courses they want to match on in the match\_profile
  - Added more preferences to Match\_Profile table to allow for more added to the scoring algorithm
  - Added a Match\_Profile\_Course table for many-to-many mapping between users and courses they want to match on
  - GetStudyBuddyMatches (Procedure)
    - Reduces candidate set via WITH (CTEs)
      - Only allows matches for users who share at least one course
      - candidate\_base aggregates shared course count and other data
    - Score\_candidates used formula based on college, style, meeting preference, goals, time, noise, shared courses, and age proximity
  - Based on the new tables and GetStudyBuddyMatches procedure
    - The in depth work happens on a much smaller candidate pool (users who share courses instead of all users)
    - Primary key on match\_profile\_course (user\_id, course\_id) help MySQL used index based joins for user/course relationships
    - Instead of  $O(U * G)$  scans over users and groups, we now have  $O(C * \log C)$  over a candidate pool C

#### 4. Concurrency in Groups

- Original join group logic used (check count then insert)
  - Vulnerable when two users clicking at the same time could possibly get in if the study group is only one away to being full
  - JoinGroupWithLock procedure uses FOR UPDATE to lock relevant rows until transactions complete
    - This allows for the join operation to be atomic even under concurrent requests
    - Group overfill is avoided
    - Complexity checks are the same as before with O(Log N) thanks to idx\_gm\_group\_user, but with isolation

#### 5. Message and Request

- In early versions of DMs it was possible to
  - Send multiple messages without the target user accepting the message request
    - We wanted it such that a user can send one message and then the other user must accept the message\_request (sending the one message is not mandatory)
    - This lead to multiple overlapping requests between the same pair.
  - Direct\_Conversation and Direct\_Message tables were added with constraints
    - uq\_message\_request\_pair and uq\_dc\_pair to make sure the conversations and requests were unique pairs
  - Procedures:
    - StartDirectConversation
      - Normalized pairs such that u1, u2 is the same as u2, u1
      - Updated/Inserts a pending message\_request
    - GetDmInboxForUser
      - Uses idx\_dm\_convo\_time to find the latest message per conversation
      - Join Pending requests to show whether a conversation is still waiting to be accepted
    - RespondToMessageRequest
      - Accept and the conversation continues
      - Reject and the conversation is cleaned up as well as messages
- The DM system is now consent based and the accept/reject works.

- Queries use composite indexes

## Vivek Reddy Bhimavarapu

### 1. ACID CONSTRAINTS

ACID transactions were required because when creating a quiz with multiple questions and answers, if something goes wrong halfway through (like a network error), you don't want a partially created quiz. ACID transactions ensure everything succeeds together or everything fails together.

Implementation in StudyBuddy:

User submits form with: quiz title, description, questions, and answers.

`conn.start_transaction()` is used to ensure the entire quiz creation process is atomic.

Real-World Example:

If inserting quiz → success, question 1 → success, question 2 → fails, then the entire quiz is rolled back so nothing is partially saved.

### 2. Parameterized SQL Queries

SQL Injection is a security vulnerability where bad actors insert malicious SQL code into input fields.

Dangerous Example (DON'T DO THIS):

```
email = request.get_json()["email"]
```

Safe Example (WHAT WE DO):

```
email = request.get_json()["email"]
```

The database treats the input as **DATA**, not SQL code. Even if user enters ' `OR '1'='1`', it's treated as a literal string to search for.

Where We Use This:

- Login:

```
cursor.execute("SELECT ... WHERE email = %s", (email,))
```

- Quizzes:

```
cursor.execute("INSERT INTO Quiz ... VALUES (%s, %s, %s, %s)",
```

```
(title, desc, course_id, creator_id))
```

- Flashcards:  

```
cursor.execute("INSERT INTO flashcardset ... VALUES (%s, %s, %s, %s)", (title, desc, course_id, creator_id))
```
- All user data updates

### 3. Proper Error Handling

Transaction Error Handling Pattern:

```
try:  
    # db ops  
except:  
    rollback  
finally:  
    close
```

What This Does:

- **try:** Attempts database operation
- **except:** If error occurs, rolls back all changes
- **finally:** Always closes connections (prevents resource leaks)

Example Flow:

- Quiz creation starts
- Insert quiz: Success ✓
- Insert question 1: Success ✓
- Insert question 2: FAILS ✗
- Exception caught

- Rollback executed → removes quiz and question 1
- Connection closed
- Error returned to user

#### 4. Database Constraints

Foreign Key Constraints

NOT NULL Constraints

Data Type Validation:

- Points are numbers, not text
- IDs are integers
- Descriptions have max length

These prevent invalid data from entering the database and enforce referential integrity.

#### 5. Summary of Benefits

ACID transactions, parameterized queries, error handling, and constraints make StudyBuddy reliable and safe for students to use.